

Winner Plc — Testing for Salary Discrimination

Introduction

At the last board meeting of **Winner Plc**, concerns were raised that women executives were being underpaid relative to men — on average by about £8,400. This notebook investigates whether this gap is statistically significant, and whether it is explained by differences in **experience** or represents genuine **gender-based discrimination**.

The dataset contains salaries for a sample of 50 executives (24 men, 26 women).

```
In [1]: # --- Import libraries ---
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

# scipy.stats gives us statistical tests like the t-test
from scipy import stats

# statsmodels is used for regression modelling (like lm() in R)
import statsmodels.formula.api as smf
import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor

from patsy import dmatrix
import warnings
warnings.filterwarnings('ignore')

# Display settings
pd.set_option('display.float_format', '{:.4f}'.format)
pd.set_option('display.max_columns', 20)
sns.set_theme(style='whitegrid')
plt.rcParams['figure.figsize'] = (10, 5)

In [2]: # --- Load the Winner Plc data ---
winner = pd.read_csv("https://raw.githubusercontent.com/kostis-christodoulou/LBS/01_websites_github_repos/data_analytics_executives/session03/Concept_Check_2.html")

print('Shape:', winner.shape)
print('\nColumn types:')
print(winner.dtypes)
print('\nFirst few rows:')
winner.head()
```

Shape: (50, 4)

Column types:

```
employee    int64
gender       str
experience   int64
salary       int64
dtype: object
```

First few rows:

```
Out [2]:
```

	employee	gender	experience	salary
0	1	male	16	79400
1	2	male	25	67400
2	3	male	15	66500
3	4	male	33	72600
4	5	male	16	63600

1. Salary on Gender

Can we conclude that there is a **significant difference** between male and female salaries?


```
In [3]: # --- Visualise salary distributions by gender ---
# Left panel: density plot (smoothed histogram)
# Right panel: horizontal boxplot
# Remember: the line in the middle of a boxplot is the MEDIAN, not the me

fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Density plot
sns.kdeplot(
    data=winner, x='salary', hue='gender',
    fill=True, alpha=0.3, ax=axes[0]
)
axes[0].set_title('Salary Distribution by Gender – Density Plot')
axes[0].set_xlabel('Salary (£)')

# Horizontal boxplot
sns.boxplot(
    data=winner, x='salary', y='gender',
    hue='gender', palette='Set2',
    legend=False, ax=axes[1]
)
axes[1].set_title('Salary Distribution by Gender – Boxplot')
axes[1].set_xlabel('Salary (£)')
axes[1].set_ylabel('')

plt.tight_layout()
plt.show()
```



```
In [4]: # --- Summary statistics by gender ---
# Equivalent to mosaic::favstats(salary ~ gender) in R
# We also compute the standard error: SE = SD / sqrt(n)

(
  winner
  .groupby('gender')['salary']
  .agg(['count', 'min', 'median', 'mean', 'max', 'std'])
  .assign(se=lambda df: df['std'] / np.sqrt(df['count']))
  .round(2)
  .rename(columns={'count': 'n'})
)
```

```
Out [4]:
```

	n	min	median	mean	max	std	se
female	26	45600	62650.0000	62576.9200	76400	7336.7700	1438.8600
male	24	53100	72400.0000	71008.3300	82000	7235.2900	1476.9000

Confidence Intervals

We can construct CIs 'by hand' using the t-distribution critical values. The t-critical value depends on the sample size and confidence level.

```
In [6]: # --- t-critical values for manual CI construction ---
# Equivalent to qt(0.975, 23) and qt(0.975, 25) in R
# n_male = 24 → df = 23
# n_female = 26 → df = 25

for label, df in [('Male (n=24)', 23), ('Female (n=26)', 25)]:
    t95 = stats.t.ppf(0.975, df=df)
    t99 = stats.t.ppf(0.995, df=df)
    t999 = stats.t.ppf(0.9995, df=df)
    print(f'{label}:')
    print(f' 95% CI t-critical : {t95:.4f}')
    print(f' 99% CI t-critical : {t99:.4f}')
    print(f' 99.9% CI t-critical: {t999:.4f}')
    print()
```

Male (n=24):
 95% CI t-critical : 2.0687
 99% CI t-critical : 2.8073
 99.9% CI t-critical: 3.7676

Female (n=26):
 95% CI t-critical : 2.0595
 99% CI t-critical : 2.7874
 99.9% CI t-critical: 3.7251

```
In [13]: def two_sample_ttest_summary(a, b, col, label_a, label_b):
  x, y = a[col].dropna(), b[col].dropna()
  t, p = stats.ttest_ind(x, y, equal_var=False) # Welch's
  df_w = (x.var()/len(x) + y.var()/len(y))*2 / (
    (x.var()/len(x))*2/(len(x)-1) + (y.var()/len(y))*2/(len(y)-1)
  )
  diff = x.mean() - y.mean()
  se = np.sqrt(x.var()/len(x) + y.var()/len(y))
  ci = stats.t.interval(0.95, df=df_w, loc=diff, scale=se)

  print(f"--- Welch Two-Sample T-Test: {col} | {label_a} vs {label_b} ---")
  print(f" n : {len(x)} vs {len(y)}")
  print(f" mean : {x.mean():.2f} vs {y.mean():.2f}")
  print(f" std : {x.std(ddof=1):.2f} vs {y.std(ddof=1):.2f}")
  print(f" diff : {diff:.2f} ({label_a} minus {label_b})")
  print(f" 95% CI : [{ci[0]:.2f}, {ci[1]:.2f}]")
  print(f" t : {t:.4f}")
  print(f" df : {df_w:.2f}")
  print(f" p-value : {p:.6f} {'***' if p<0.001 else '**' if p<0.01 else ''}")
```

```
In [14]: two_sample_ttest_summary(
  winner.loc[winner['gender'] == 'male'],
  winner.loc[winner['gender'] == 'female'],
  col='salary',
  label_a='Male', label_b='Female'
)
```

```
--- Welch Two-Sample T-Test: salary | Male vs Female ---
n : 24 vs 26
mean : 71,008.33 vs 62,576.92
std : 7,235.29 vs 7,336.77
diff : 8,431.41 (Male minus Female)
95% CI : [4,285.13, 12,577.69]
t : 4.0891
df : 47.78
p-value : 0.000165 ***
```

```
In [21]: def fit_and_plot(formula, data, model_name="Model", x_var=None):
  model = smf.ols(formula=formula, data=data).fit()

  # — Compact summary —
  rse = np.sqrt(model.mse_resid)
  print(f"\n{'='*55}")
  print(f" {model_name}")
  print(f"{'='*55}")
  print(f" R²: {model.rsquared:.4f} Adj R²: {model.rsquared_adj:.4f}")
  print(f" Observations: {int(model.nobs)} df residual: {int(model.df_resid)}")
  print(f"{'-'*55}")
  print(f" {'Coefficient':<20} {'Estimate':>10} {'Std Err':>10} {'t':>10}")
  print(f"{'-'*55}")
```

```

for name, coef, se, tval, pval in zip(
    model.params.index,
    model.params,
    model.bse,
    model.tvalues,
    model.pvalues
):
    sig = "***" if pval < 0.001 else "**" if pval < 0.01 else "*" if
    print(f" {name:<20} {coef:>10.4f} {se:>10.4f} {tval:>8.3f} {pval}
print(f"{' '*55}\n")

if x_var is None:
    x_var = formula.split('~')[1].strip().split('+')[0].strip()

y_var      = formula.split('~')[0].strip()
residuals = model.resid
is_cat     = data[x_var].dtype == object or str(data[x_var].dtype) ==

fig = plt.figure(figsize=(15, 10))
gs  = fig.add_gridspec(2, 2, height_ratios=[1.6, 1], hspace=0.35, wsp
ax_main = fig.add_subplot(gs[0, :])
ax_res  = fig.add_subplot(gs[1, 0])
ax_qq   = fig.add_subplot(gs[1, 1])

# — 1. Actual vs Predicted —————
if is_cat:
    sns.stripplot(x=data[x_var], y=data[y_var],
                  alpha=0.3, color='black', s=4, ax=ax_main, label='A
    group_means = data.groupby(x_var)[y_var].mean()
    ax_main.hlines(group_means.values,
                   xmin=[i - 0.3 for i in range(len(group_means))],
                   xmax=[i + 0.3 for i in range(len(group_means))],
                   colors='red', linewidth=2, label='Predicted mean')
else:
    ax_main.scatter(data[x_var], data[y_var],
                   alpha=0.3, color='black', s=15, label='Actual')
    ax_main.scatter(data[x_var], model.fittedvalues,
                   alpha=0.3, color='red', s=15, label='Predicted')

ax_main.set_title(f'Actual vs Predicted – x: {x_var}', fontsize=13)
ax_main.set_xlabel(x_var)
ax_main.set_ylabel(y_var)
ax_main.legend()

# — 2. Residuals vs Fitted —————
ax_res.scatter(model.fittedvalues, residuals,
               alpha=0.3, color='steelblue', s=15)
ax_res.axhline(0, color='red', linewidth=1, linestyle='--')
ax_res.set_title('Residuals vs Fitted', fontsize=11)
ax_res.set_xlabel('Fitted Values')
ax_res.set_ylabel('Residuals')

# — 3. Q-Q plot —————
(osm, osr), (slope, intercept, _) = stats.probplot(residuals, dist='n
ax_qq.scatter(osm, osr, alpha=0.3, color='steelblue', s=15)
ax_qq.plot(osm, slope * np.array(osm) + intercept,
           color='red', linewidth=1, linestyle='--')
ax_qq.set_title('Q-Q Plot of Residuals', fontsize=11)
ax_qq.set_xlabel('Theoretical Quantiles')
ax_qq.set_ylabel('Sample Quantiles')

```

```
# --- Suptitle with all three stats ---
plt.suptitle(
    f'{model_name} | R² = {model.rsquared:.3f} | Adj R² = {model.
    fontsize=13, fontweight='bold', y=1.01
)
plt.tight_layout()
plt.show()

return model
```

```
In [22]: # --- Regression: salary ~ gender ---
# The slope on 'genderMale' is the mean salary difference between men and
# Compare it with the observed difference above – they should be identical

model1 = fit_and_plot('salary ~ gender',
                      data=winner,
                      model_name='Salary ~ Gender',
                      x_var='gender')
```

=====

Salary ~ Gender




=====

R²: 0.2581 Adj R²: 0.2427 Residual SE: 7288.3247
 Observations: 50 df residual: 48

Coefficient	Estimate	Std Err	t	P> t
Intercept	62576.9231	1429.3581	43.780	0.0000 ***
gender[T.male]	8431.4103	2063.1007	4.087	0.0002 ***

=====

Salary ~ Gender | R² = 0.258 | Adj R² = 0.243 | Residual SE = 7288.325



2. Experience on Gender

At the board meeting, someone suggested that the salary gap might reflect **differences in experience** rather than discrimination. Men in the sample have ~21 years experience on average vs ~7 years for women.


```
In [9]: # --- Visualise experience distributions by gender ---

fig, axes = plt.subplots(1, 2, figsize=(14, 5))

sns.kdeplot(
    data=winner, x='experience', hue='gender',
    fill=True, alpha=0.3, ax=axes[0]
)
axes[0].set_title('Experience Distribution by Gender – Density Plot')
axes[0].set_xlabel('Years of Experience')

sns.boxplot(
    data=winner, x='experience', y='gender',
    hue='gender', palette='Set2',
    legend=False, ax=axes[1]
)
axes[1].set_title('Experience Distribution by Gender – Boxplot')
axes[1].set_xlabel('Years of Experience')
axes[1].set_ylabel('')

plt.tight_layout()
plt.show()
```



```
In [18]: # --- Summary statistics: experience by gender ---

(
    winner
    .groupby('gender')['experience']
    .agg(['count', 'min', 'median', 'mean', 'max', 'std'])
    .assign(se=lambda df: df['std'] / np.sqrt(df['count']))
    .round(2)
    .rename(columns={'count': 'n'})
)
```

```
Out[18]:
```

	n	min	median	mean	max	std	se
gender							
female	26	0	3.0000	7.3800	29	8.5100	1.6700
male	24	1	19.5000	21.1200	44	10.9200	2.2300

```
In [16]: # --- t-test: is the experience gap between genders significant? ---
# If experience differs significantly AND experience drives salary,
# the salary gap may be explained by experience rather than discrimination

two_sample_ttest_summary(
    winner.loc[winner['gender'] == 'male'],
    winner.loc[winner['gender'] == 'female'],
    col='experience',
    label_a='Male', label_b='Female'
)

--- Welch Two-Sample T-Test: experience | Male vs Female ---
n      : 24 vs 26
mean   : 21.12 vs 7.38
std    : 10.92 vs 8.51
diff   : 13.74 (Male minus Female)
95% CI : [8.13, 19.35]
t      : 4.9348
df     : 43.47
p-value : 0.000012 ***
```

3. Salary on Experience

We now examine the direct relationship between salary and experience.

```
In [18]: # --- Correlation matrix: salary vs experience ---

fig, ax = plt.subplots(figsize=(5, 4))

sns.heatmap(
    winner[['experience', 'salary']].corr(),
    annot=True, fmt='.3f',
    cmap='coolwarm', center=0, vmin=-1, vmax=1,
    square=True, linewidths=0.5,
    ax=ax
)

ax.set_title('Correlation: Experience vs Salary')
plt.tight_layout()
plt.show()
```



```
In [23]: # --- Model 2: salary ~ experience ---
# A simple linear regression of salary on years of experience

model2 = fit_and_plot('salary ~ experience',
                      data=winner,
                      model_name='Salary ~ Experience',
                      x_var='experience')
```

=====


Salary ~ Experience

=====

R²: 0.6451 Adj R²: 0.6377 Residual SE: 5041.3443
 Observations: 50 df residual: 48

Coefficient	Estimate	Std Err	t	P> t
Intercept	58704.7613	1107.8104	52.992	0.0000 ***
experience	566.4691	60.6512	9.340	0.0000 ***

=====

Salary ~ Experience | $R^2 = 0.645$ | Adj $R^2 = 0.638$ | Residual SE = 5041.344

4. Salary on Gender + Experience


Extend the model to include both predictors. Does gender remain significant once we control for experience?

```
In [24]: # --- Model 3: salary ~ experience + gender ---

model3 = fit_and_plot('salary ~ experience + gender',
                      data=winner,
                      model_name='Salary ~ Experience',
                      x_var='experience')

=====
Salary ~ Experience
=====
R2: 0.6474    Adj R2: 0.6324    Residual SE: 5078.0653
Observations: 50    df residual: 47
```

Coefficient	Estimate	Std Err	t	P> t
Intercept	58573.9755	1140.4683	51.360	0.0000 ***
gender[T.male]	983.2175	1770.7599	0.555	0.5814
experience	542.0658	75.2593	7.203	0.0000 ***

Salary ~ Experience | $R^2 = 0.647$ | Adj $R^2 = 0.632$ | Residual SE = 5078.065

5. Improving the Model — Non-linear Experience Effect

Someone at the meeting argued that salary increases **faster** in early career years than in later years. We check whether a square-root transformation of experience fits better.

```
In [27]: # --- Scatterplots with trend lines to visualise the shape of the relation

fig, axes = plt.subplots(1, 3, figsize=(18, 5))

# Panel 1: linear trend (OLS)
sns.regplot(
    data=winner, x='experience', y='salary',
    ax=axes[0], scatter_kws={'alpha': 0.5}
)
axes[0].set_title('salary ~ experience (linear)')

# Panel 2: colour by gender with linear fit per group
for gender, grp in winner.groupby('gender'):
    sns.regplot(
        data=grp, x='experience', y='salary',
        ax=axes[1], label=gender,
        scatter_kws={'alpha': 0.5}
    )
axes[1].set_title('salary ~ experience (by gender)')
axes[1].legend()
```

```
# Panel 3: LOWESS smooth (non-parametric) to let the data suggest the shape
sns.regplot(
    data=winner, x='experience', y='salary',
    ax=axes[2],
    lowess=True, # equivalent to geom_smooth(se=FALSE) in R
    scatter_kws={'alpha': 0.5},
    line_kws={'color': 'red'}
)
axes[2].set_title('salary ~ experience (LOWESS smooth)')

plt.tight_layout()
plt.show()
print('The LOWESS curve suggests salary grows quickly early on, then flat
```



The LOWESS curve suggests salary grows quickly early on, then flattens – a `sqrt()` shape.

```
In [25]: # --- Model 4: salary ~ sqrt(experience) + gender ---
# A sqrt() transformation lets the marginal effect of experience decrease

winner_aug = winner.assign(sqrt_experience=lambda df: np.sqrt(df['experie


model4 = fit_and_plot('salary ~ sqrt_experience + C(gender)',
    data=winner_aug,
    model_name='Salary ~ Experience',
    x_var='sqrt_experience')
```

```
=====
Salary ~ Experience
=====
```

```
R2: 0.7403    Adj R2: 0.7292    Residual SE: 4358.0788
Observations: 50    df residual: 47
```

Coefficient	Estimate	Std Err	t	P> t
Intercept	54759.0620	1196.2514	45.776	0.0000 ***
C(gender) [T.male]	-73.4243	1533.2672	-0.048	0.9620
sqrt_experience	3713.7379	397.5891	9.341	0.0000 ***

```
=====
```

Salary ~ Experience | $R^2 = 0.740$ | Adj $R^2 = 0.729$ | Residual SE = 4358.079

In [26]: # --- Model 5: salary ~ sqrt(experience) – remove non-significant gender


```
model5 = fit_and_plot('salary ~ sqrt_experience',
                      data=winner_aug,
                      model_name='Salary ~ Experience',
                      x_var='sqrt_experience')
```

=====
Salary ~ Experience
=====

R^2 : 0.7403 Adj R^2 : 0.7348 Residual SE: 4312.5484
Observations: 50 df residual: 48

Coefficient	Estimate	Std Err	t	P> t
Intercept	54760.0484	1183.5782	46.267	0.0000 ***
sqrt_experience	3702.4314	316.5508	11.696	0.0000 ***

=====

Salary ~ Experience | $R^2 = 0.740$ | Adj $R^2 = 0.735$ | Residual SE = 4312.548

In [39]: # --- Model comparison table (equivalent to huxtable::huxreg) ---

```
models = {
  'M1: ~gender':      model1,
  'M2: ~exp':         model2,
  'M3: ~exp+gender':  model3,
  'M4: ~sqrt+gender': model4,
  'M5: ~sqrt':        model5,
}

comparison = pd.DataFrame({
  name: {
    'N':      int(m.nobs),
    'R2':    round(m.rsquared, 4),
    'Adj. R2': round(m.rsquared_adj, 4),
    'Residual SE': round(np.sqrt(m.mse_resid), 0),
    'AIC':     round(m.aic, 1),
  }
  for name, m in models.items()
})

print('Model Comparison:')
comparison
```

Model Comparison:

Out [39]:

	M1: ~gender	M2: ~exp	M3: ~exp+gender	M4: ~sqrt+gender	M5: ~sqrt
N	50.0000	50.0000	50.0000	50.0000	50.0000
R²	0.6451	0.6474	0.6474	0.7403	0.7403
Adj. R²	0.6377	0.6324	0.6324	0.7292	0.7348
Residual SE	5041.0000	5078.0000	5078.0000	4358.0000	4313.0000
AIC	996.4000	998.1000	998.1000	982.8000	980.8000

Conclusions

1. The raw salary gap (~£8,400) **is statistically significant** when gender alone is the predictor.
2. However, men in this sample have on average **~14 more years of experience** than women.
3. Once experience is included in the model, **gender is no longer statistically significant**.
4. The best-fitting model uses **$\sqrt{(\text{experience})}$** as the predictor (Adj. R² is highest), confirming that salary growth decelerates with tenure.
5. **No evidence of discrimination** in this data, once experience differences are accounted for-- but it's a cheap argument